

# Live Code Execution Backend

## A Scalable Architecture for Remote Code Execution

Nguyen Bao Duy

Edtronaut - SWE Intern Case Study

January 22, 2026

# Outline

1 Introduction

2 System Architecture

3 Core Technical Design

4 Design Trade-offs

5 Reliability & Scalability

6 Conclusion

# Problem Statement

## The Challenge

Build a backend service that allows users to write and execute code in multiple programming languages remotely.

## Key Requirements:

- ✓ Support multiple languages (JavaScript, Python, Java)
- ✓ Handle concurrent users safely
- ✓ Protect against malicious code (timeouts, resource limits)
- ✓ Provide session management with autosave
- ✓ Return execution results asynchronously

## Core Technical Challenge

How do we execute untrusted user code safely, reliably, and at scale?



# Project Overview

## What Was Built:

- RESTful API for code sessions
- Queue-based async execution
- Multi-language support
- Rate limiting & spam protection
- Session state management

## Technology Stack:

- **Runtime:** Node.js + TypeScript
- **Queue:** Redis + BullMQ
- **Database:** SQLite
- **Architecture:** API + Worker pattern


## Design Philosophy

Prioritize **simplicity** and **reliability** over premature optimization.

# Outline

- 1 Introduction
- 2 System Architecture
- 3 Core Technical Design
- 4 Design Trade-offs
- 5 Reliability & Scalability
- 6 Conclusion

# High-Level Architecture



## Three-Component Design

- **API Server**: Handles HTTP requests, manages sessions, enqueues jobs
- **Redis Queue**: Decouples API from execution, enables async processing
- **Worker Process**: Executes code in isolated environment, updates results

# Outline

- 1 Introduction
- 2 System Architecture
- 3 Core Technical Design
- 4 Design Trade-offs
- 5 Reliability & Scalability
- 6 Conclusion

# Design Decision: Async Execution with Queue

## Problem

Running user code takes time (1-5 seconds). If API waits, it blocks other requests and degrades performance.

## Solution: Queue-Based Architecture

- API adds job to Redis queue and returns **immediately**
- Separate Worker process dequeues and executes code
- Results written to database for later retrieval

## Benefits:

- ✓ API stays fast and responsive
- ✓ Independent scaling of API and Workers
- ✓ Automatic retries on failure

## Trade-offs:

- ✗ Increased system complexity
- ✗ Results not instant (polling required)
- ✗ Redis dependency

# State Management: Execution Lifecycle

## State Ownership

Clear separation: **API creates, Worker executes**

State	Set by	Responsibility
QUEUED	API	Creates the execution record
RUNNING	Worker	Picks up the job and starts it
COMPLETED / FAILED / TIMEOUT	Worker	Writes final result to database

## Database as Source of Truth

SQLite stores authoritative state.

# Code Isolation: Sandboxing Strategy

## Problem

User code is **untrusted**. Must prevent infinite loops, excessive output, and file/network access.

## Solution: Child Process Isolation

- Run code in separate child process (not main server)
- **5-second timeout:** Kill process if execution exceeds limit
- **1MB output limit:** Monitor stdout/stderr in real-time, kill if exceeded
- Process killed immediately with SIGKILL signal

## Trade-off

**Production:** Use Docker or gVisor for stronger security.

# Autosave Spam Protection

## Problem

Users typing rapidly generate **100+ requests/minute**, blocking SQLite.

## Solution: Adaptive Throttling

### Three-layer strategy:

- ① **Throttling:** Max 1 write/second per session
- ② **Coalescing:** Only save latest code
- ③ **Forced write:** If pending > 5s, force write

### Guarantees:

- ✓ Latest code saved
- ✓ No delay > 5s

### Trade-off:

- ✗ In-memory state
- ✗ Lost on crash

# Execution Environment: Resource Limits

## Protection Mechanisms

Multiple layers of resource enforcement:

Resource	Limit	Action
Execution time	5 seconds	Kill → TIMEOUT
Output size	1 MB	Kill → Error
Concurrent jobs	5 jobs	Queue
Executions/session	5/minute	HTTP 429
Cooldown	2 seconds	Prevent re-run

## Defense in Depth

Multiple limits ensure **no single user** can overwhelm the system.

# Outline

- 1 Introduction
- 2 System Architecture
- 3 Core Technical Design
- 4 Design Trade-offs
- 5 Reliability & Scalability
- 6 Conclusion

# Trade-off Analysis Framework

## Decision-Making Approach

Every architectural choice involves trade-offs.

**Framework:** Problem → Decision → Benefits vs. Costs

## Optimization Priorities

- ① **Simplicity:** Single database, minimal orchestration
- ② **Reliability:** Automatic retries, graceful failure handling
- ③ **Speed:** Async queue, non-blocking API
- ④ **Safety:** Timeouts, output limits, rate limiting

## Not Optimized For

- High availability (single SQLite = single point of failure)
- Horizontal API scaling (SQLite doesn't support concurrent writes well)

# Trade-off #1: SQLite vs PostgreSQL

## SQLite (Chosen)

### Why:

- ✓ Zero setup, file-based
- ✓ Perfect for development
- ✓ Simple deployment

### Cost:

- ✗ Single writer only
- ✗ No concurrent API instances
- ✗ Not production-ready

## PostgreSQL (Alternative)

### Benefits:

- ✓ Concurrent writes
- ✓ Read replicas
- ✓ Connection pooling

### Cost:

- ✗ Complex setup
- ✗ Requires hosting
- ✗ Migration overhead

# Trade-off #2: Child Process vs Docker

## Child Process (Chosen)

### Why:

- ✓ Simple, no dependencies
- ✓ Fast startup (<50ms)
- ✓ Easy local development

### Cost:

- ✗ Limited isolation
- ✗ OS-level limits only
- ✗ Potential security risks

## Docker (Alternative)

### Benefits:

- ✓ Full sandboxing
- ✓ Network isolation
- ✓ No host access

### Cost:

- ✗ Slower startup (~500ms)
- ✗ Docker daemon required
- ✗ More complex error handling

# Trade-off #3: Polling vs WebSocket/SSE

## Polling (Chosen)

### Why:

- ✓ Simple client code
- ✓ No connection state
- ✓ Stateless API design
- ✓ Works with load balancers

### Cost:

- ✗ Higher latency (poll interval)
- ✗ Unnecessary requests
- ✗ More server load

## WebSocket/SSE (Alternative)

### Benefits:

- ✓ Real-time updates
- ✓ Low latency
- ✓ Better UX

### Cost:

- ✗ Connection management
- ✗ Reconnection logic
- ✗ Scaling complexity
- ✗ Stateful API required

# Outline

- 1 Introduction
- 2 System Architecture
- 3 Core Technical Design
- 4 Design Trade-offs
- 5 Reliability & Scalability
- 6 Conclusion

# Idempotency: Four-Layer Protection

## Problem

Users may click "Run" multiple times quickly.

System must prevent duplicate executions without rejecting legitimate requests.

## Solution: Defense in Depth

### Layer 1: Active Execution Check

Before creating job, check if session has QUEUED or RUNNING execution.

### Layer 2: Cooldown Period

2-second cooldown after each execution. Prevents rapid-fire requests.

### Layer 3: Rate Limit

Maximum 5 executions per minute per session. Protects against abuse.

### Layer 4: Database Constraint

SQLite unique constraint on execution ID as last-resort safety net.

# Failure Handling: Stalled Job Detection

## Problem

Worker may crash mid-execution. How does system recover?

## BullMQ Automatic Recovery

- ① Worker marks job as RUNNING and acquires lock (30s duration)
- ② If worker crashes, lock expires
- ③ BullMQ detects stalled job via stalledInterval check (30s)
- ④ Job automatically retried (up to 3 attempts)
- ⑤ After exhausting retries → failed event → DB updated to FAILED

## Guarantees:

- ✓ No jobs stuck forever
- ✓ Automatic retry on crash

## Out of Scope:

- ✗ True exactly-once execution
- ✗ Dead letter queue (DLQ)

# Horizontal Scaling Strategy

## Stateless API Design

No per-request or per-user state in API memory.

All persistent state in shared storage: Database + Redis (queues, rate limits).

### Scaling API Servers:

- Any API instance handles any request
- Load balancer distributes traffic
- **Bottleneck:** Database write contention
- **Solution:** PostgreSQL with row-level locking & connection pooling

### Scaling Workers:

- Workers pull jobs from Redis queue
- Each worker: 5 concurrent jobs
- Independent processes (deploy, scale, restart separately)
- **Throughput:** Linear scaling (until CPU/Redis bottleneck)

# Scalability Bottlenecks & Mitigations

Bottleneck	Current Mitigation	Production Solution
SQLite write conflicts	Autosave throttling, single DB	PostgreSQL with connection pooling
Redis memory growth	Auto-cleanup (1h, 1000 jobs)	Max queue size limit, priority queues
Worker CPU overload	Concurrency limit (5 jobs)	Kubernetes auto-scaling, resource quotas
Long execution times	5-second global timeout	Per-language timeouts, adaptive limits
Excessive output	1MB limit, kill process	Stream output, early truncation

## Design for Current Needs, Plan for Growth

Current architecture handles **moderate load** (100s of users).

Clear migration path documented for production scale (1000s of users).

# Outline

- 1 Introduction
- 2 System Architecture
- 3 Core Technical Design
- 4 Design Trade-offs
- 5 Reliability & Scalability
- 6 Conclusion

# What Was Achieved

## Successfully Implemented

- **✓ Async queue-based execution** with Redis + BullIMQ
- **✓ Multi-language support** (JavaScript, Python, Java)
- **✓ Session management** with autosave and throttling
- **✓ Code isolation** with timeout and output limits
- **✓ Idempotency protection** with 4-layer defense
- **✓ Automatic failure recovery** via BullIMQ retry logic
- **✓ Scalable worker architecture** with horizontal scaling

# Production Readiness Gaps

## What Would Be Improved with More Time

### 1. Stronger Container Isolation

Replace child processes with Docker or gVisor for full sandboxing

### 2. PostgreSQL Migration

Enable concurrent writes, read replicas, and horizontal API scaling

### 3. Observability

Structured logging, Prometheus metrics, BullMQ dashboard, alerting

### 4. Real-time Updates

WebSocket/SSE for pushing execution status instead of polling

### 5. Advanced Queue Management

Max queue size, priority queues, per-language isolation, backpressure

Thank You

# Live Code Execution Backend

A Case Study in Scalable System Design

**GitHub:** [github.com/BuhDuy256/live-code-execution-backend](https://github.com/BuhDuy256/live-code-execution-backend)

**Tech Stack:** Node.js, TypeScript, Redis, BullMQ, SQLite

*Questions?*